

# Exploratory Data Analysis

# Introduction

- Applying data mining (InfoVis as well) techniques requires gaining useful insights into the input data first
  - We saw this in the previous lecture
- Exploratory Data Analysis (EDA) helps to achieve this
- EDA offers several techniques to comprehend data
- But EDA is more than a library of data analysis techniques
- EDA is an approach to data analysis
- EDA involves inspecting data without any assumptions
  - Mostly using information graphics
  - Modern InfoVis tools use many of the EDA techniques which we study later
- Insights gained from EDA help selecting appropriate data mining (InfoVis) technique.

# Descriptive Statistics

- Descriptive statistical methods quantitatively describe the main features of data
- Main data features
  - measures of central tendency - represent a 'center' around which measurements are distributed
    - e.g. mean and median
  - measures of variability - represent the 'spread' of the data from the 'center'
    - e.g. standard deviation
  - measures of relative standing - represent the 'relative position' of specific measurements in the data
    - e.g. quantiles

# Mean

- Sum all the numbers and divide by their count

$$\bar{x} = (x_1 + x_2 + \dots + x_n) / n$$


- For the example data

- Mean =  $(2+3+4+5+6)/5$

- = 4

- 4 is the 'center'


- The information graphic used here is called a dot diagram




# Median

- The exact middle value
- When count is odd just find the middle value of the sorted data
- When count is even find the mean of the middle two values
- For example data 1
  - Median is 4
  - 4 is the 'center'
- For example data 2
  - Median is  $(3+4)/2 = 3.5$
  - 3.5 is the 'center'

Data 1




Data 2



# Median VS Mean


- When data has outliers median is more robust
  - The blue data point is the outlier in data 2
- When data distribution is skewed median is more meaningful
- For example data 1
  - Mean=4 and median=4
- For example data 2
  - Mean=24/5 and median=4



# Standard Deviation

- Computation steps
  - Compute mean
  - Compute each measurement's deviations from the mean
  - Square the deviations
  - Sum the squared deviations
  - Divide by (count-1)
  - Compute the square root

$$\sigma = \sqrt{(\sum(x_i - \bar{x})^2)/(n-1)}$$



Mean = 4

Deviations: -2, -1, 0, 1, 2


Squared deviations: 4, 1, 0, 1, 4

Sum = 10

Standard deviation =  $\sqrt{(10/4)} = 1.58$

# Quartiles

- Median is the 2<sup>nd</sup> quartile
- 1<sup>st</sup> quartile is the measurement with 25% measurements smaller and 75% larger - lower quartile (Q1)
- 3<sup>rd</sup> quartile is the measurement with 75% measurements smaller and 25% larger - upper quartile (Q3)
- Inter quartile range (IQR) is the difference between Q3 and Q1
  - $Q3 - Q1$





# Stem and Leaf Plot

- This plot organizes data for easy visual inspection
  - Min and max values
  - Data distribution
- Unlike descriptive statistics, this plot shows all the data
  - No information loss
  - Individual values can be inspected
- Structure of the plot
  - Stem - the digits in the largest place (e.g. tens place)
  - Leaves - the digits in the smallest place (e.g. ones place)
  - Leaves are listed to the left of stem separated by '|'
- Possible to place leaves from another data set to the right of the stem for comparing two data distributions

## Data

29, 44, 12, 53, 21, 34, 39, 25,  
48, 23, 17, 24, 27, 32, 34, 15,  
42, 21, 28, 37

## Stem and Leaf Plot

1 | 2 7 5

2 | 9 1 5 3 4 7 1 8

3 | 4 9 2 4 7

4 | 4 8 2


5 | 3

# Histogram/Bar Chart

- Graphical display of frequency distribution
  - Counts of data falling in various ranges (bins)
  - Histogram for numeric data
  - Bar chart for nominal data
- Bin size selection is important
  - Too small - may show spurious patterns
  - Too large - may hide important patterns
- Several Variations possible
  - Plot relative frequencies instead of raw frequencies
  - Make the height of the histogram equal to the 'relative frequency/width'
    - Area under the histogram is 1
- When observations come from continuous scale histograms can be approximated by continuous curves


## Data

29, 44, 12, 53, 21, 34, 39, 25,  
48, 23, 17, 24, 27, 32, 34, 15,  
42, 21, 28, 37



# Normal Distribution

- Distributions of several data sets are bell shaped
  - Symmetric distribution
  - With peak of the bell at the mean,  $\mu$  of the data
  - With spread (extent) of the bell defined by the standard deviation,  $\sigma$  of the data
- For example, height, weight and IQ scores are normally distributed
- **The 68-95-99.7% Rule**
  - 68% of measurements fall within  $\mu - \sigma$  and  $\mu + \sigma$
  - 95% of measurements fall within  $\mu - 2\sigma$  and  $\mu + 2\sigma$
  - 99.7% of observations fall within  $\mu - 3\sigma$  and  $\mu + 3\sigma$



# Standardization

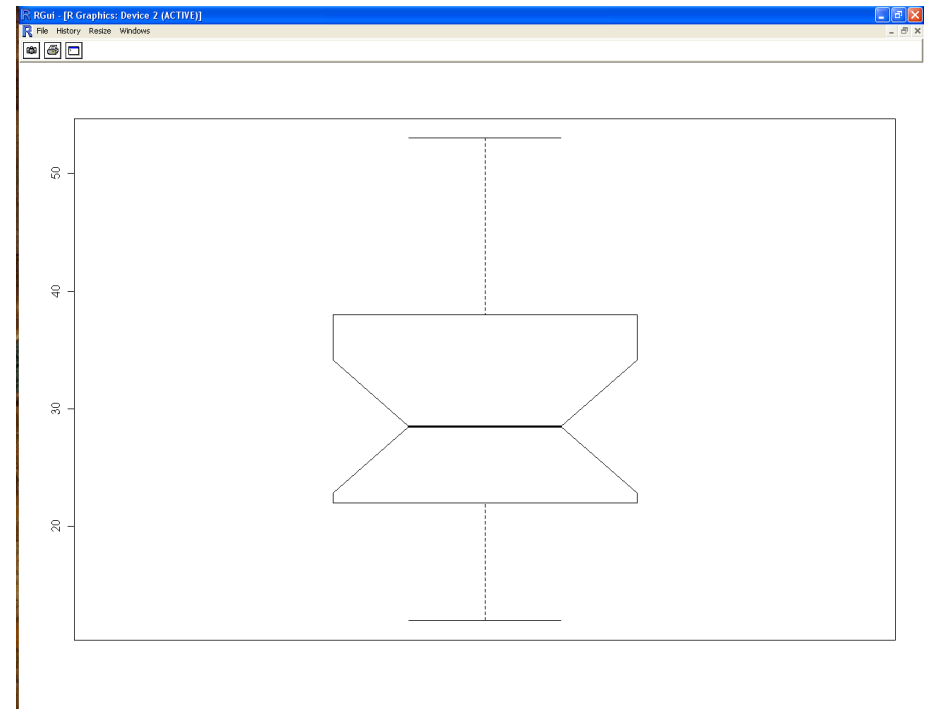
- Data sets originate from several sources and there are bound to be differences in measurements
  - Comparing data from different distributions is hard
- Standard deviation of a data set is used as a yardstick for adjusting for such distribution specific differences
- Individual measurements are converted into what are called standard measurements called z scores
- An individual measurement is expressed in terms of the number of standard deviations,  $\sigma$  it is away from the mean,  $\mu$
- Z score of  $x = (x - \mu) / \sigma$ 
  - Formula for standardizing attribute values
- Z scores are more meaningful for comparison
- When different attributes use different ranges of values, we use standardization

# Box Plot

- A five value summary plot of data
  - Minimum, maximum
  - Median
  - 1<sup>st</sup> and 3<sup>rd</sup> quartiles
- Often used in conjunction with a histogram in EDA
- Structure of the plot
  - Box represents the IQR (the middle 50% values)
  - The horizontal line in the box shows the median
  - Vertical lines extend above and below the box
  - Ends of vertical lines called whiskers indicate the max and min values
    - If max and min fall within  $1.5 \times \text{IQR}$
  - Shows outliers above/below the whiskers

## Data

29, 44, 12, 53, 21, 34, 39, 25,  
48, 23, 17, 24, 27, 32, 34, 15,  
42, 21, 28, 37



# Iris Sample Data Set


- Many of the exploratory data techniques are illustrated with the Iris Plant data set.
  - Can be obtained from the UCI Machine Learning Repository  
<http://www.ics.uci.edu/~mlearn/MLRepository.html>
  - From the statistician Douglas Fisher
  - Three flower types (classes):
    - Setosa
    - Virginica
    - Versicolour
  - Four (non-class) attributes
    - Sepal width and length
    - Petal width and length



Virginica. Robert H. Mohlenbrock.  
USDA NRCS. 1995. Northeast  
wetland flora: Field office guide to  
plant species. Northeast National  
Technical Center, Chester, PA.  
Courtesy of USDA NRCS Wetland


# Scatter Plot

- Scatter plots are two dimensional graphs with
  - explanatory attribute plotted on the x-axis
  - Response attribute plotted on the y-axis
- Useful for understanding the relationship between two attributes
- Features of the relationship
  - strength
  - shape (linear or curve)
  - Direction
  - Outliers
- Scatter plot of `iris$Petal.Width` against `iris$Petal.Length` (refer to practical 1 about IRIS data) is shown here



# Scatter Plot Matrix

- When multiple attributes need to be visualized all at once
  - Scatter plots are drawn for every pair of attributes and arranged into a 2D matrix.
- Useful for spotting relationships among attributes
  - Similar to a scatter plot
- Scatter plot matrix of IRIS data is shown here
  - Attributes are shown on the diagonal
- Later in the course we learn to use parallel coordinates for plotting multi-attribute data







## Visualization Techniques: Parallel Coordinates

- Parallel Coordinates

- Used to plot the attribute values of high-dimensional data
- Instead of using perpendicular axes, use a set of parallel axes
- The attribute values of each object are plotted as a point on each corresponding coordinate axis and the points are connected by a line
- Thus, each object is represented as a line
- Often, the lines representing a distinct class of objects group together, at least for some attributes

\_\_\_\_\_

# Parallel Coordinates Plots for Iris Data




# Visualization Techniques:


## Matrix Plots

- Matrix plots
  - Can plot the data matrix
  - This can be useful when objects are sorted according to class
  - Typically, the attributes are normalized to prevent one attribute from dominating the plot
  - Plots of similarity or distance matrices can also be useful for visualizing the relationships between objects
  - Examples of matrix plots are presented on the next two slides

# Visualization of the Iris Data Matrix



# Visualization of the Iris Correlation Matrix



# EDA Answers Questions

- All the techniques presented so far are the tools useful for EDA
- But without an understanding built from the EDA, effective use of tools is not possible
  - A detective investigating a crime scene needs tools for obtaining finger prints.
  - Also needs an understanding (common sense) to know where to look for finger prints
    - Door knobs better places than door hinges?
- EDA helps to answer a lot of questions
  - What is a typical value?
  - What is the uncertainty of a typical value?
  - What is a good distributional fit for the data?
  - What are the relationships between two attributes?
  - etc

# Next Three Lectures

- Classification
- Dr. Chenghua Lin

# Acknowledgement

- Some of the slides are based on the course slides provided by
  - Tan, Steinbach and Kumar (Introduction to Data Mining)
- Some pictures are taken from various online resources.